



# Advanced OpenCL by Example



Justin Hensley, Derek Gerstmann &  
Takahiro Harada



# Course Agenda

OpenCL Review

OpenCL Development Tips

OpenGL/OpenCL Interoperability

Rigid body particle simulation

*<15 minute Break>*

Bullet cloth

Galaxy n-body simulation

Visualization with OpenCL

# Slides from tutorial and course

<http://sa10.idav.ucdavis.edu>

# GPU Basics

**Architectures meant for throughput, not for single thread**

Not complexity of out-of-order issue

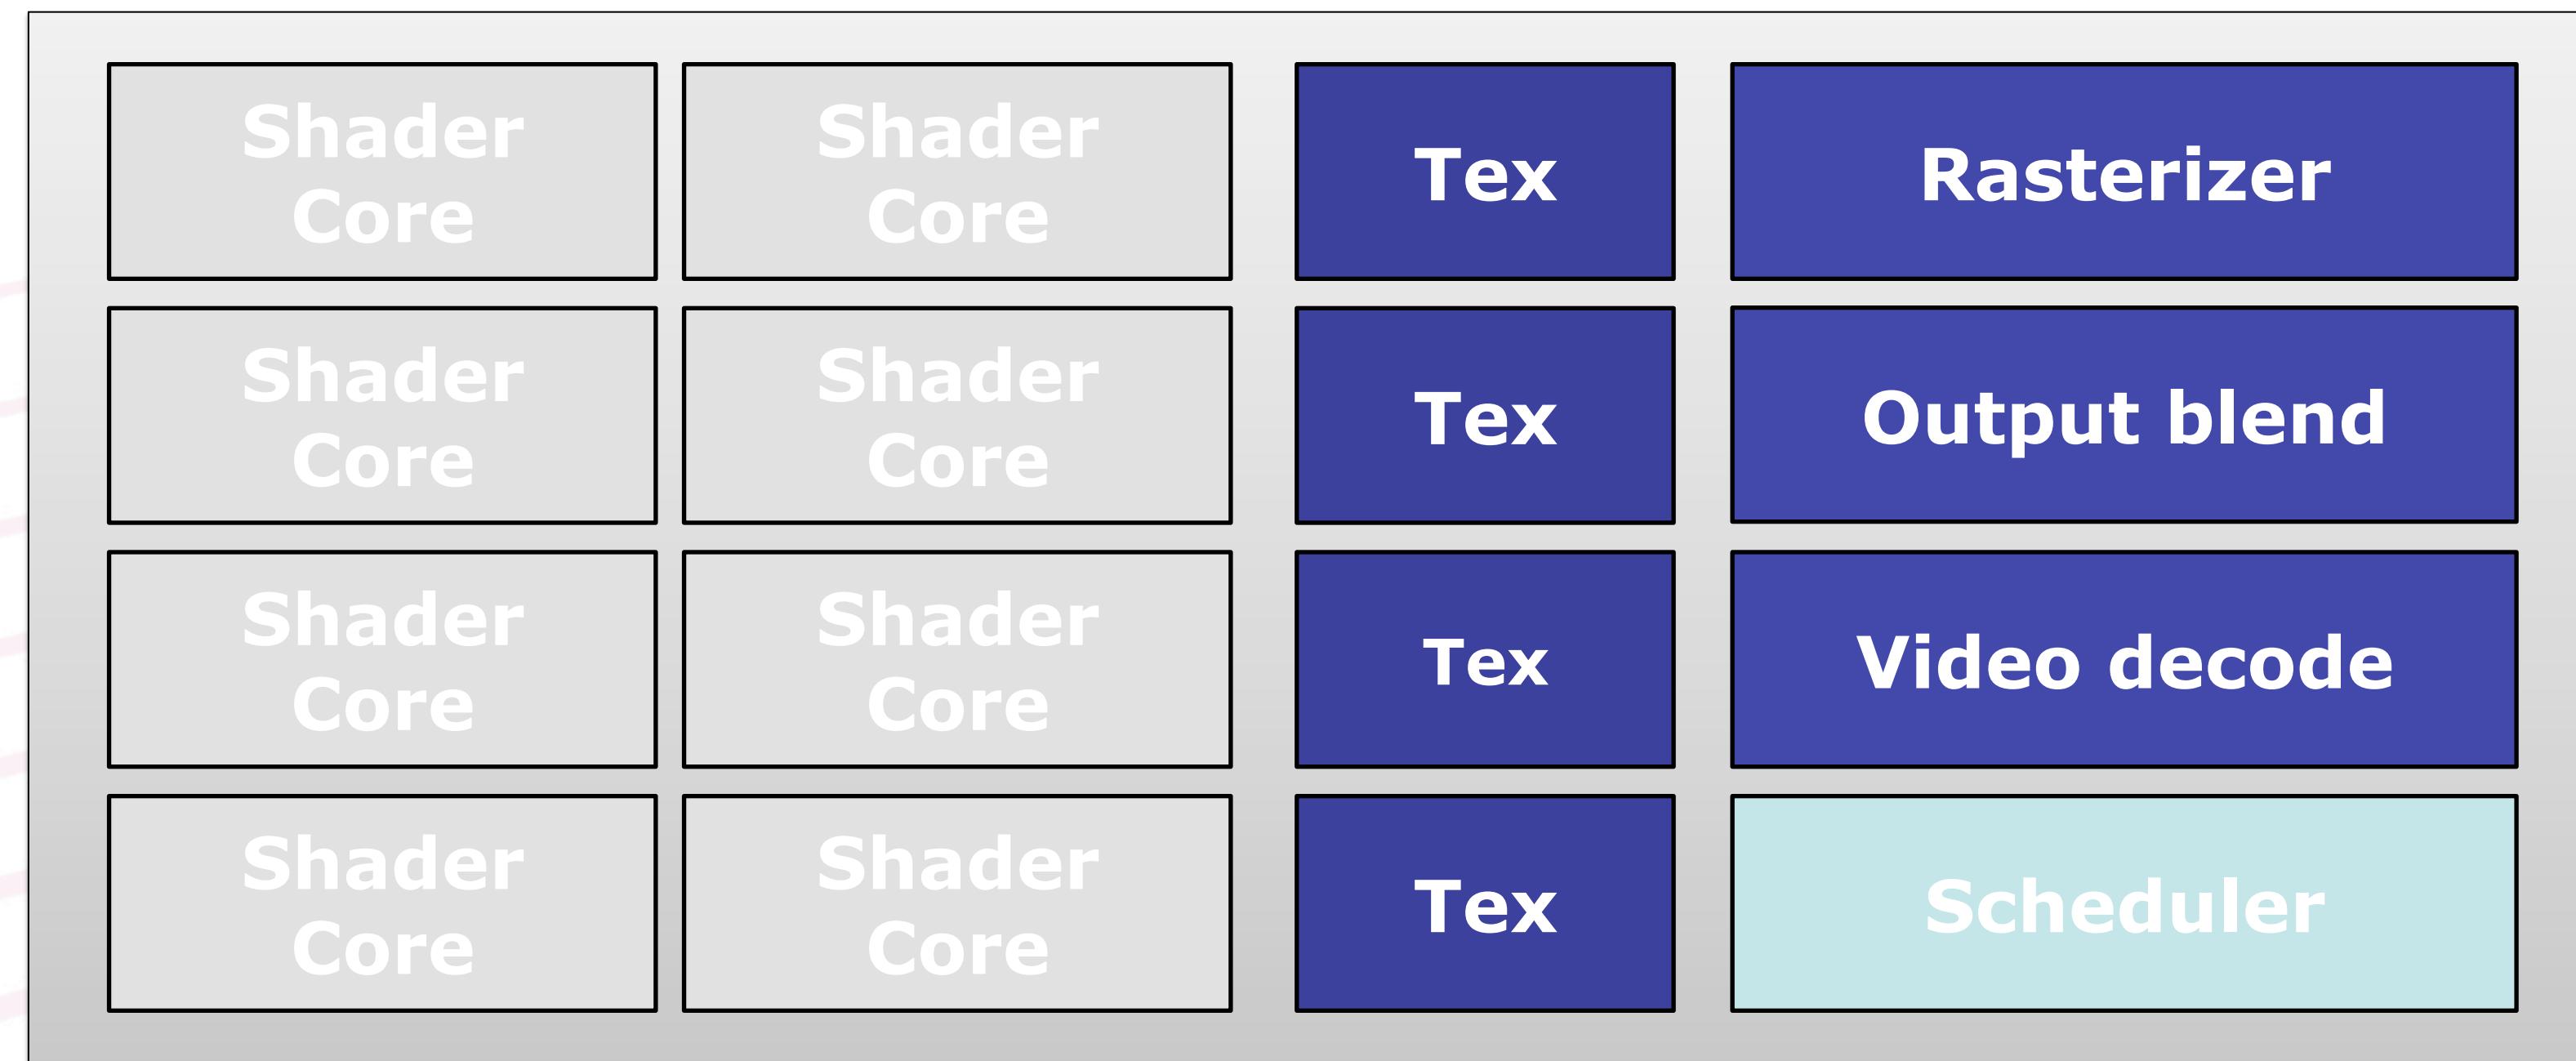
Nor of multiple issue per core

**Rendering is latency tolerant**

**Graphics generates many parallelizable tasks**

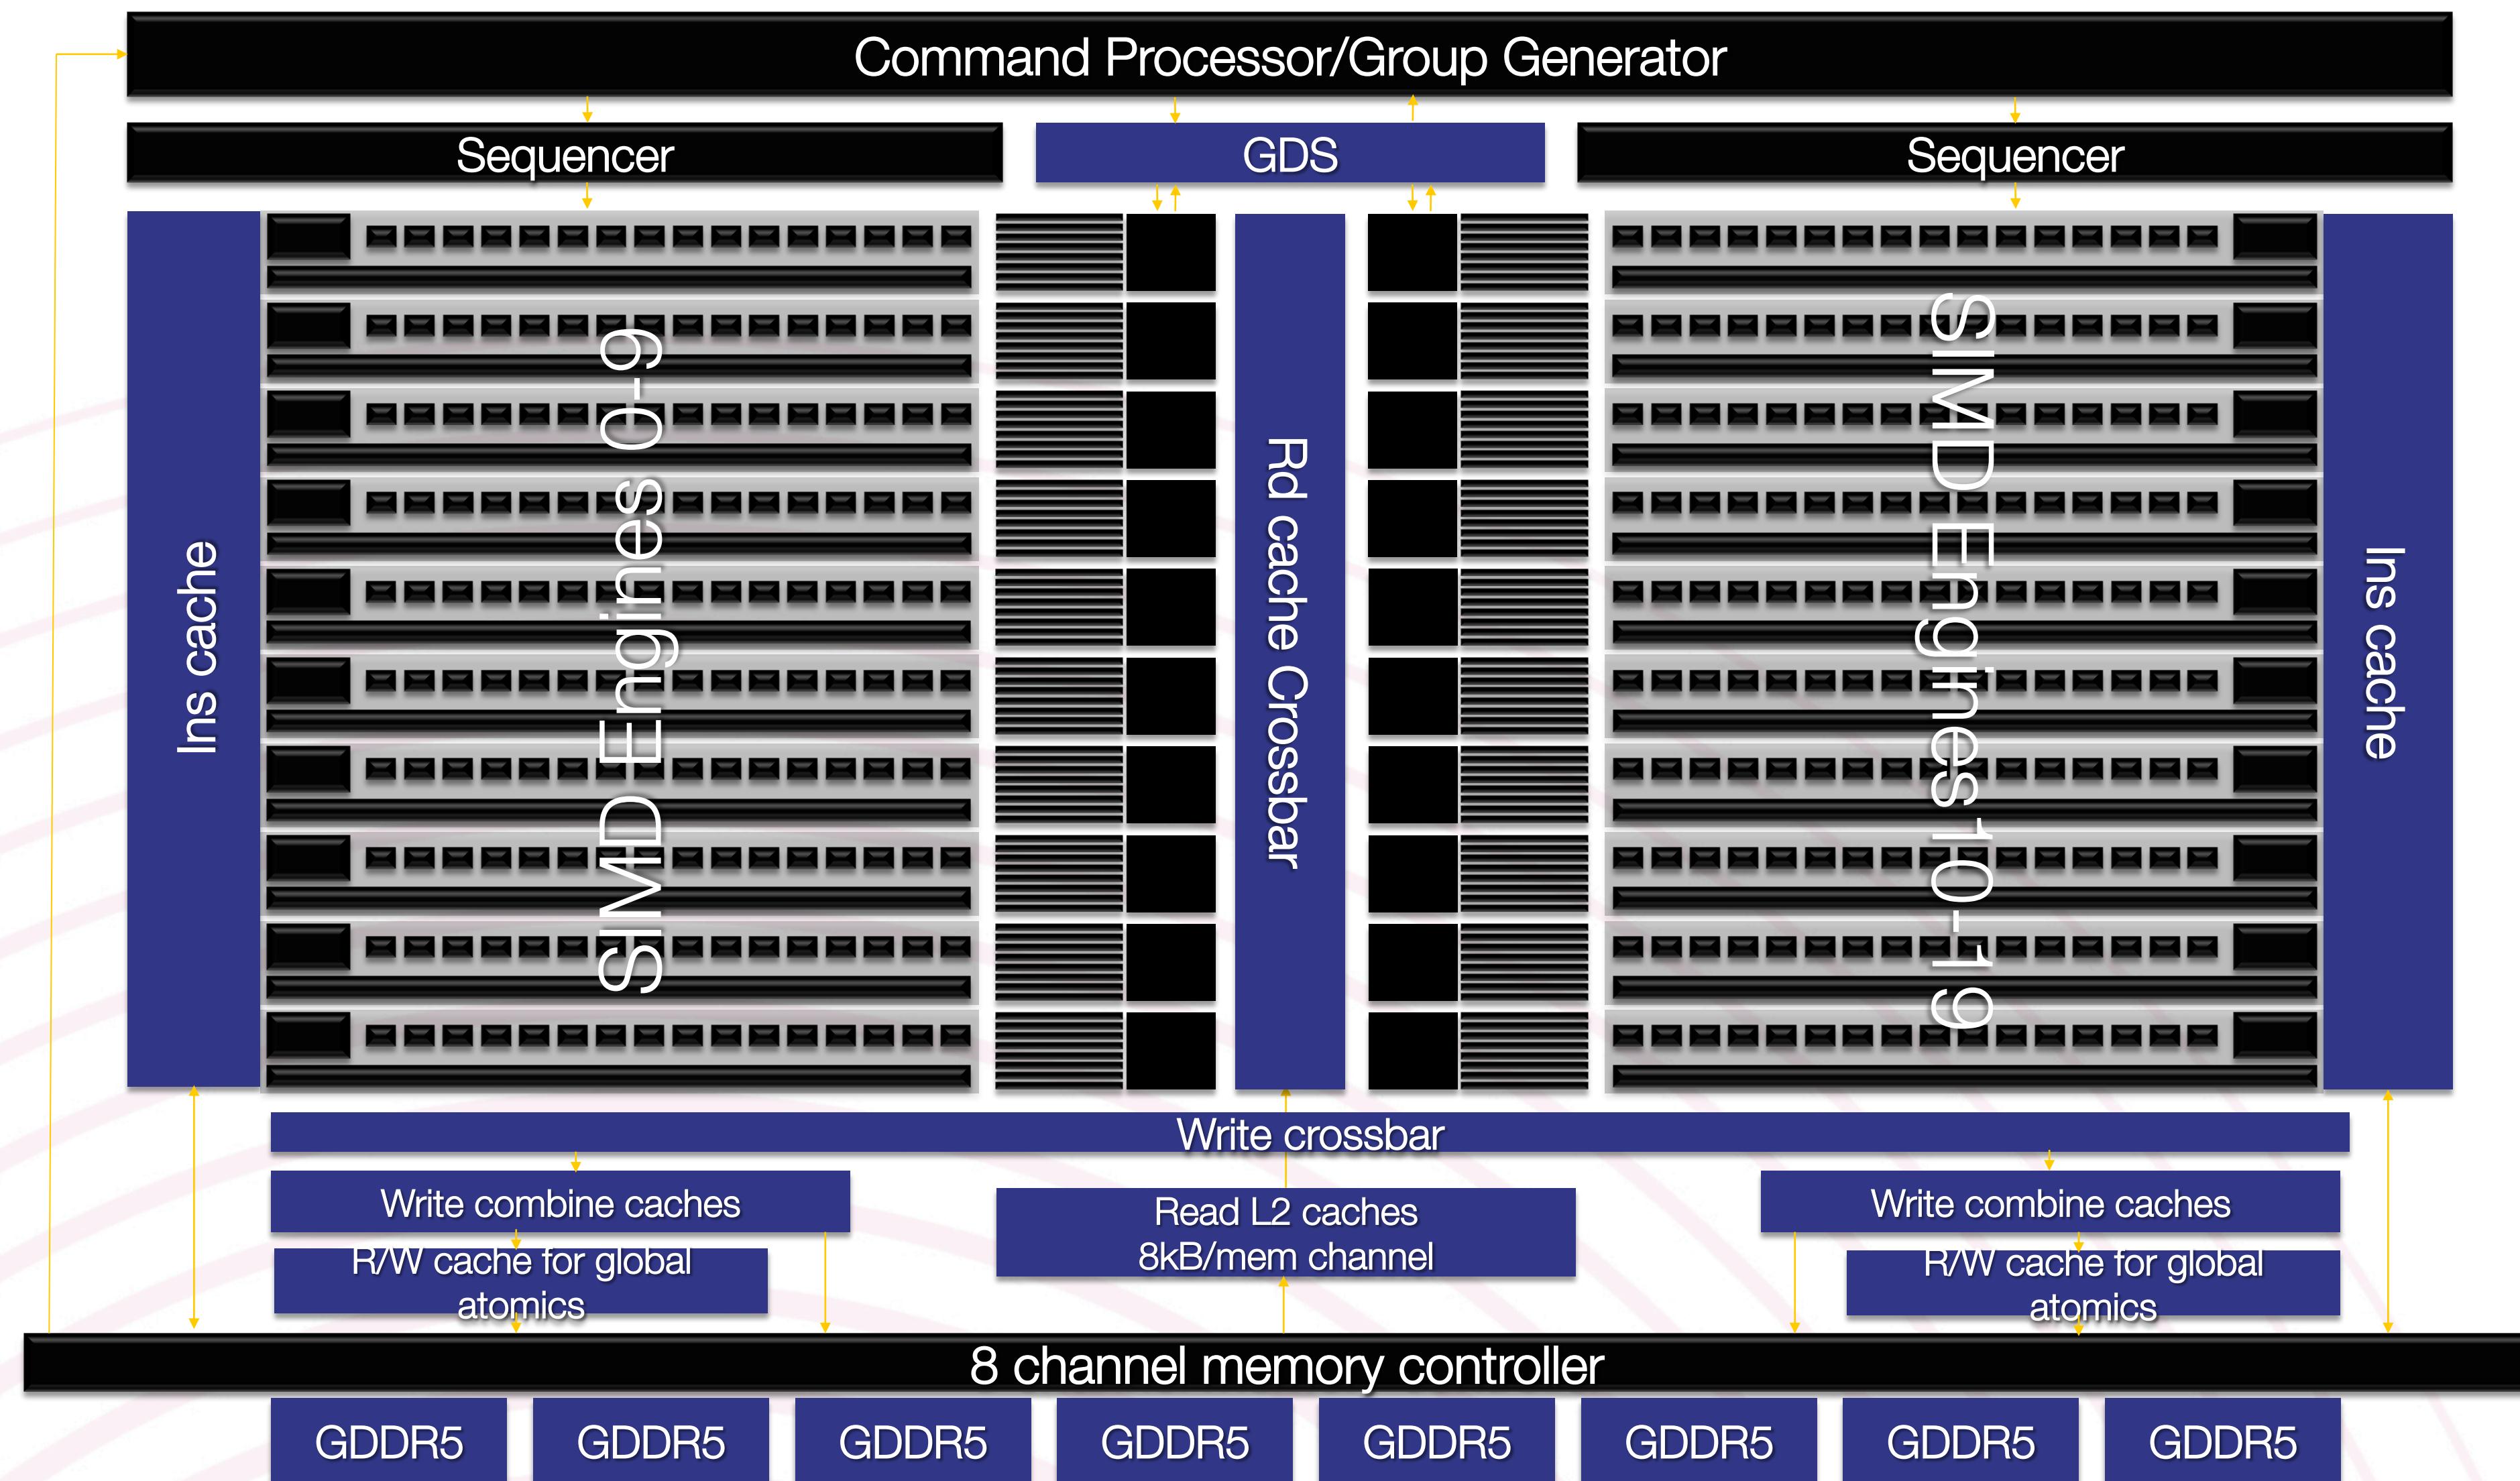
Therefore much easier to scale

# GPU In a nutshell

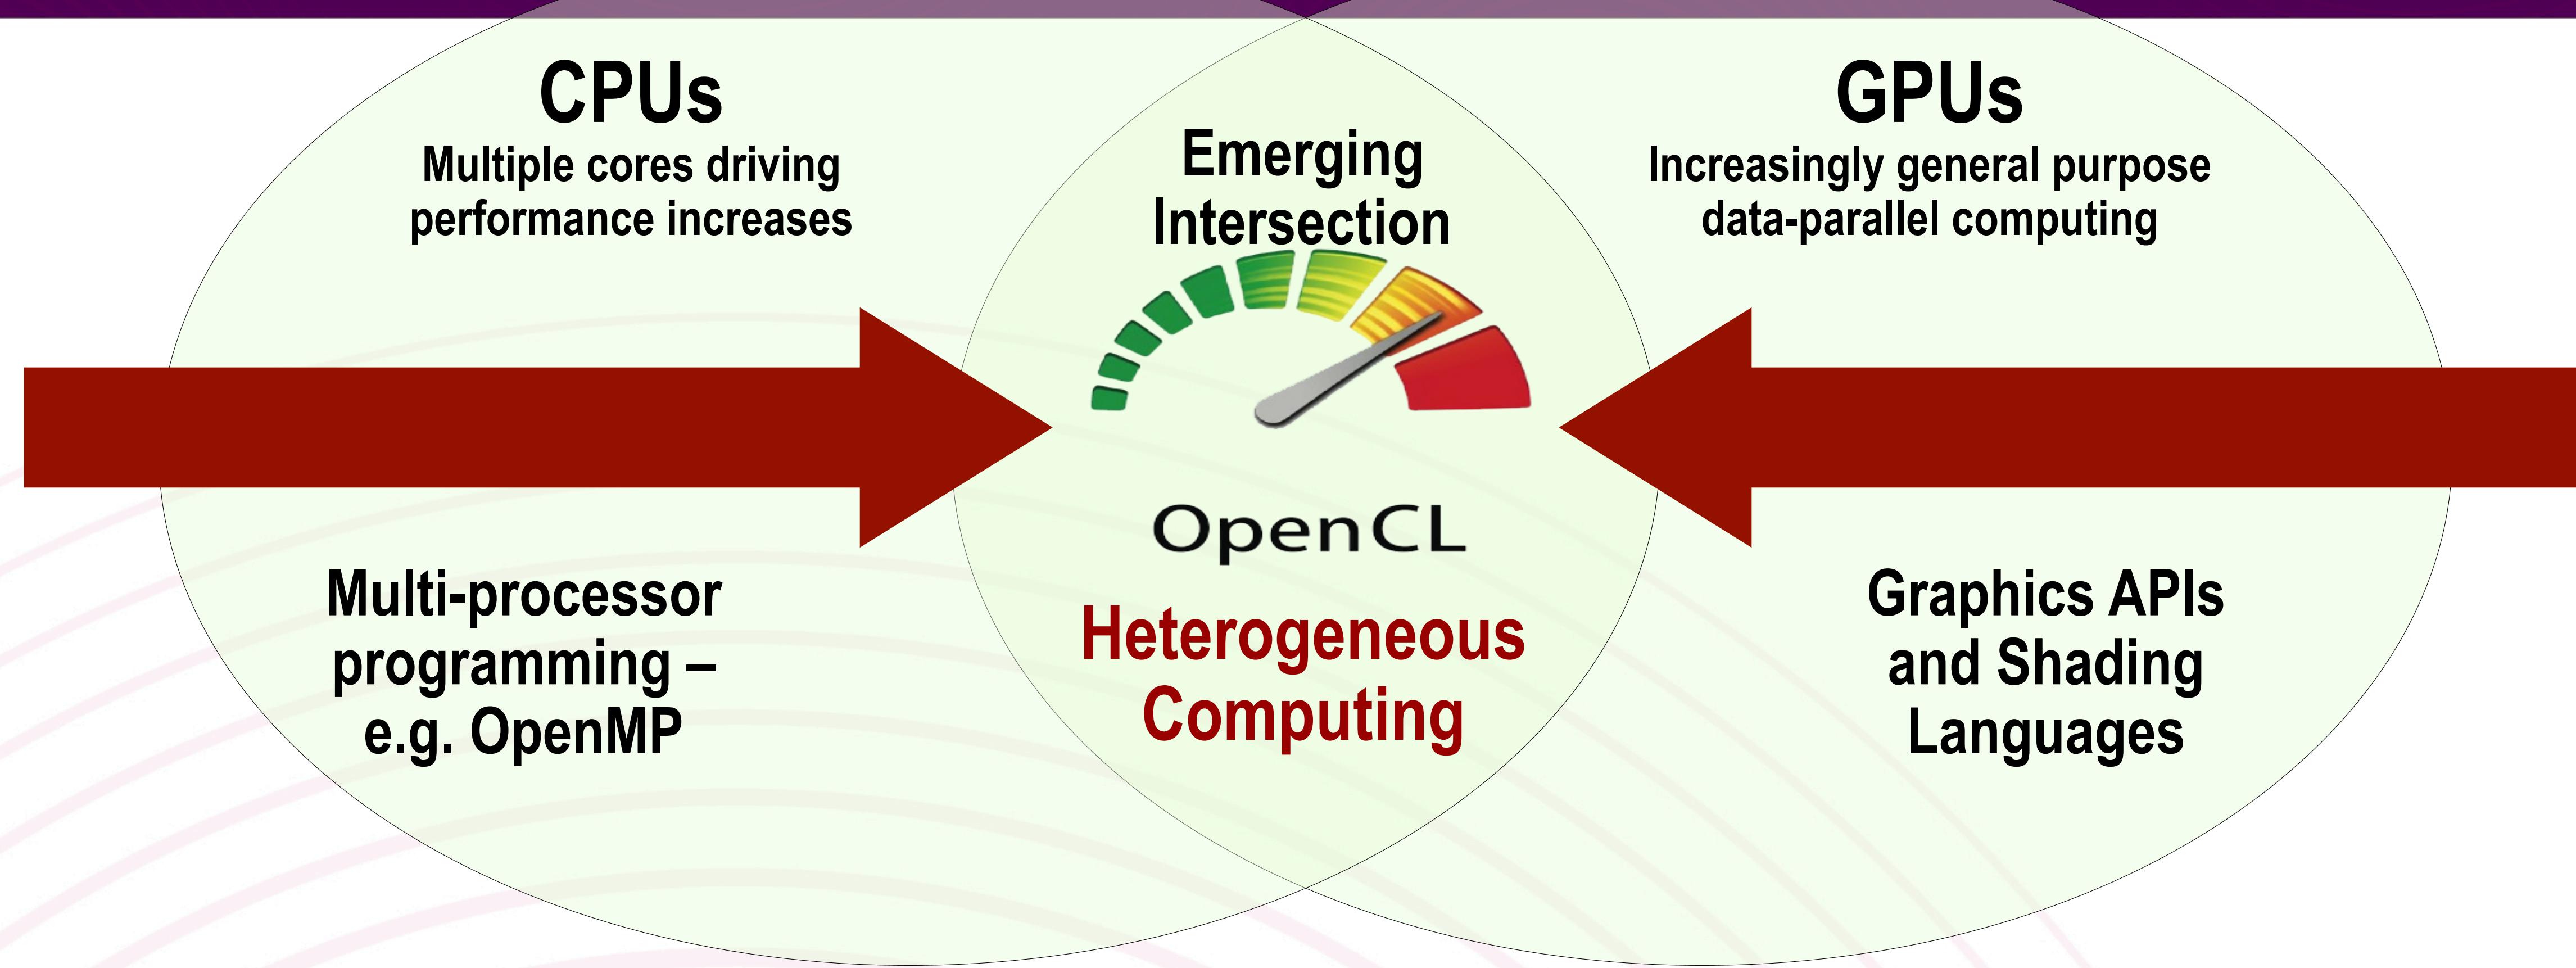


The GPU is a multicore processor optimized for graphics workloads

# Modern AMD GPU



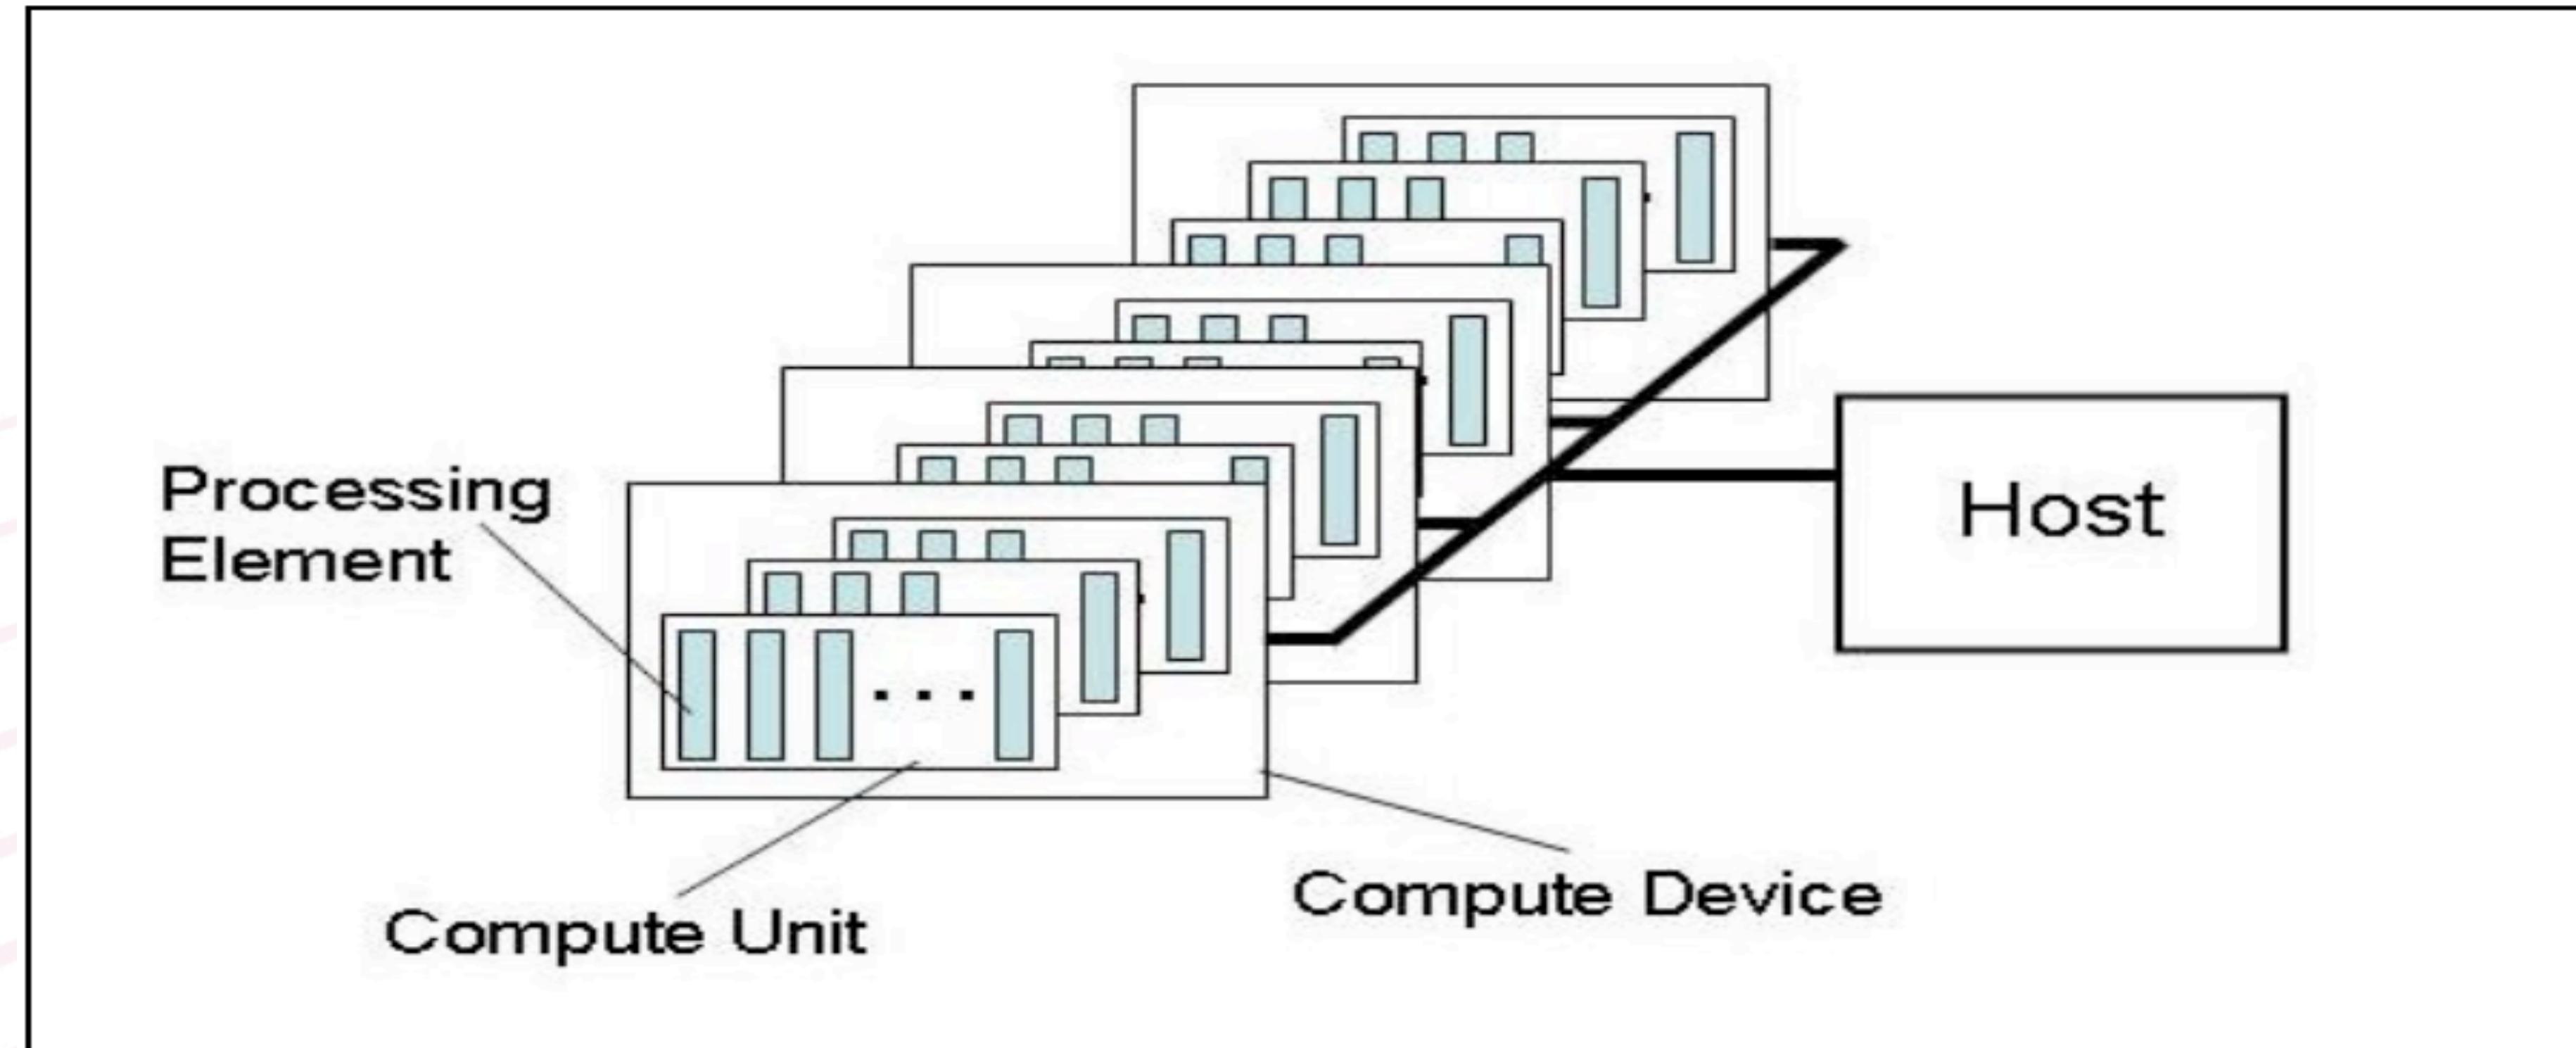
# Industry Standard for Programming Heterogeneous Platforms



## OpenCL – Open Computing Language

Open, royalty-free standard for portable, parallel programming of heterogeneous parallel computing CPUs, GPUs, and other processors

# OpenCL Platform Model

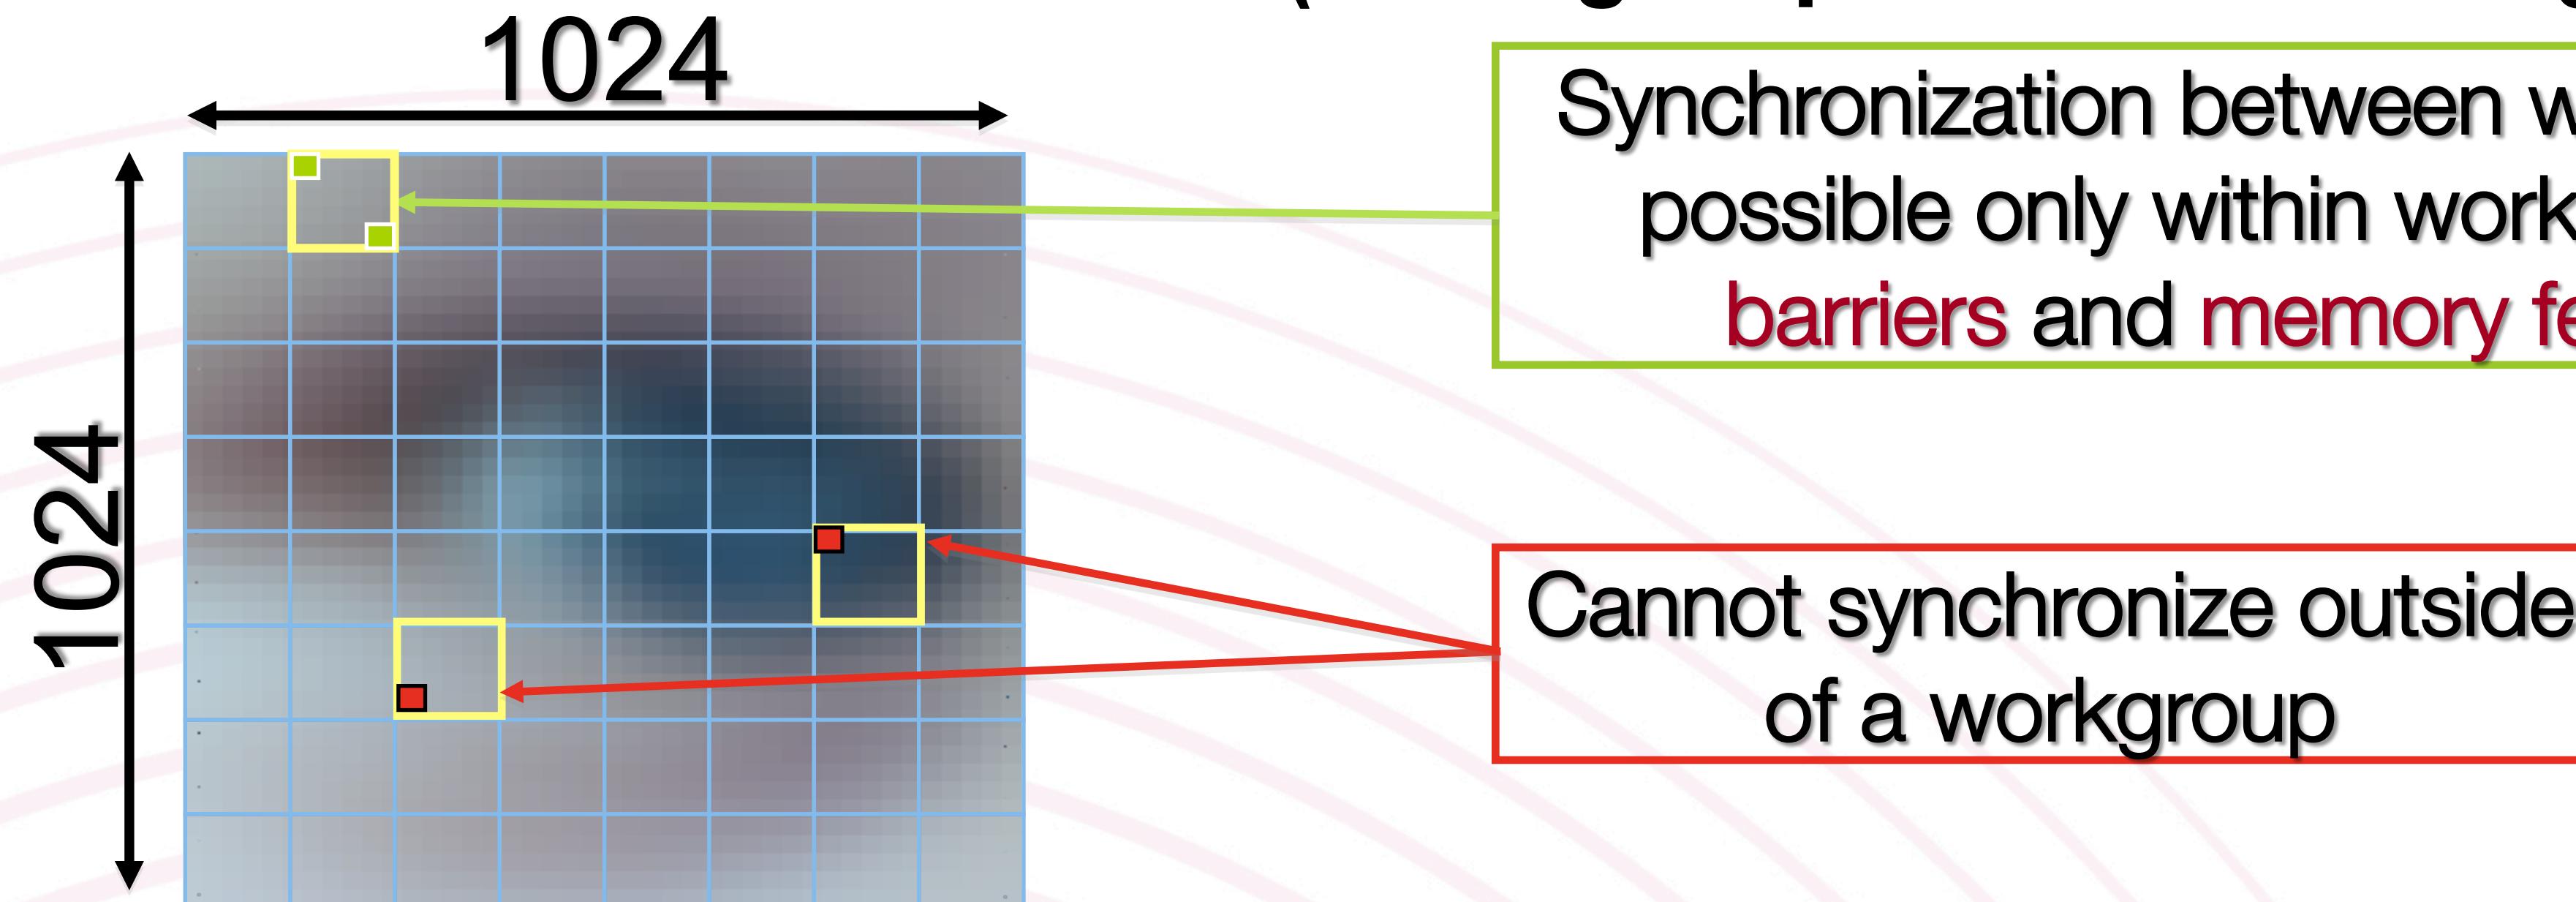


- One Host + one or more Compute Devices
  - Each Compute Device is composed of one or more Compute Units
    - Each Compute Unit is further divided into one or more Processing Elements

Each Compute Unit is further divided into one or more Processing Elements

# An N-dimension domain of work-items

- Global Dimensions:  $1024 \times 1024$  (whole problem space)
- Local Dimensions:  $128 \times 128$  (work group ... executes together)



- Choose the dimensions that are “best” for your algorithm

# OpenCL Memory Model

- **Private Memory**

- Per work-item

- **Local Memory**

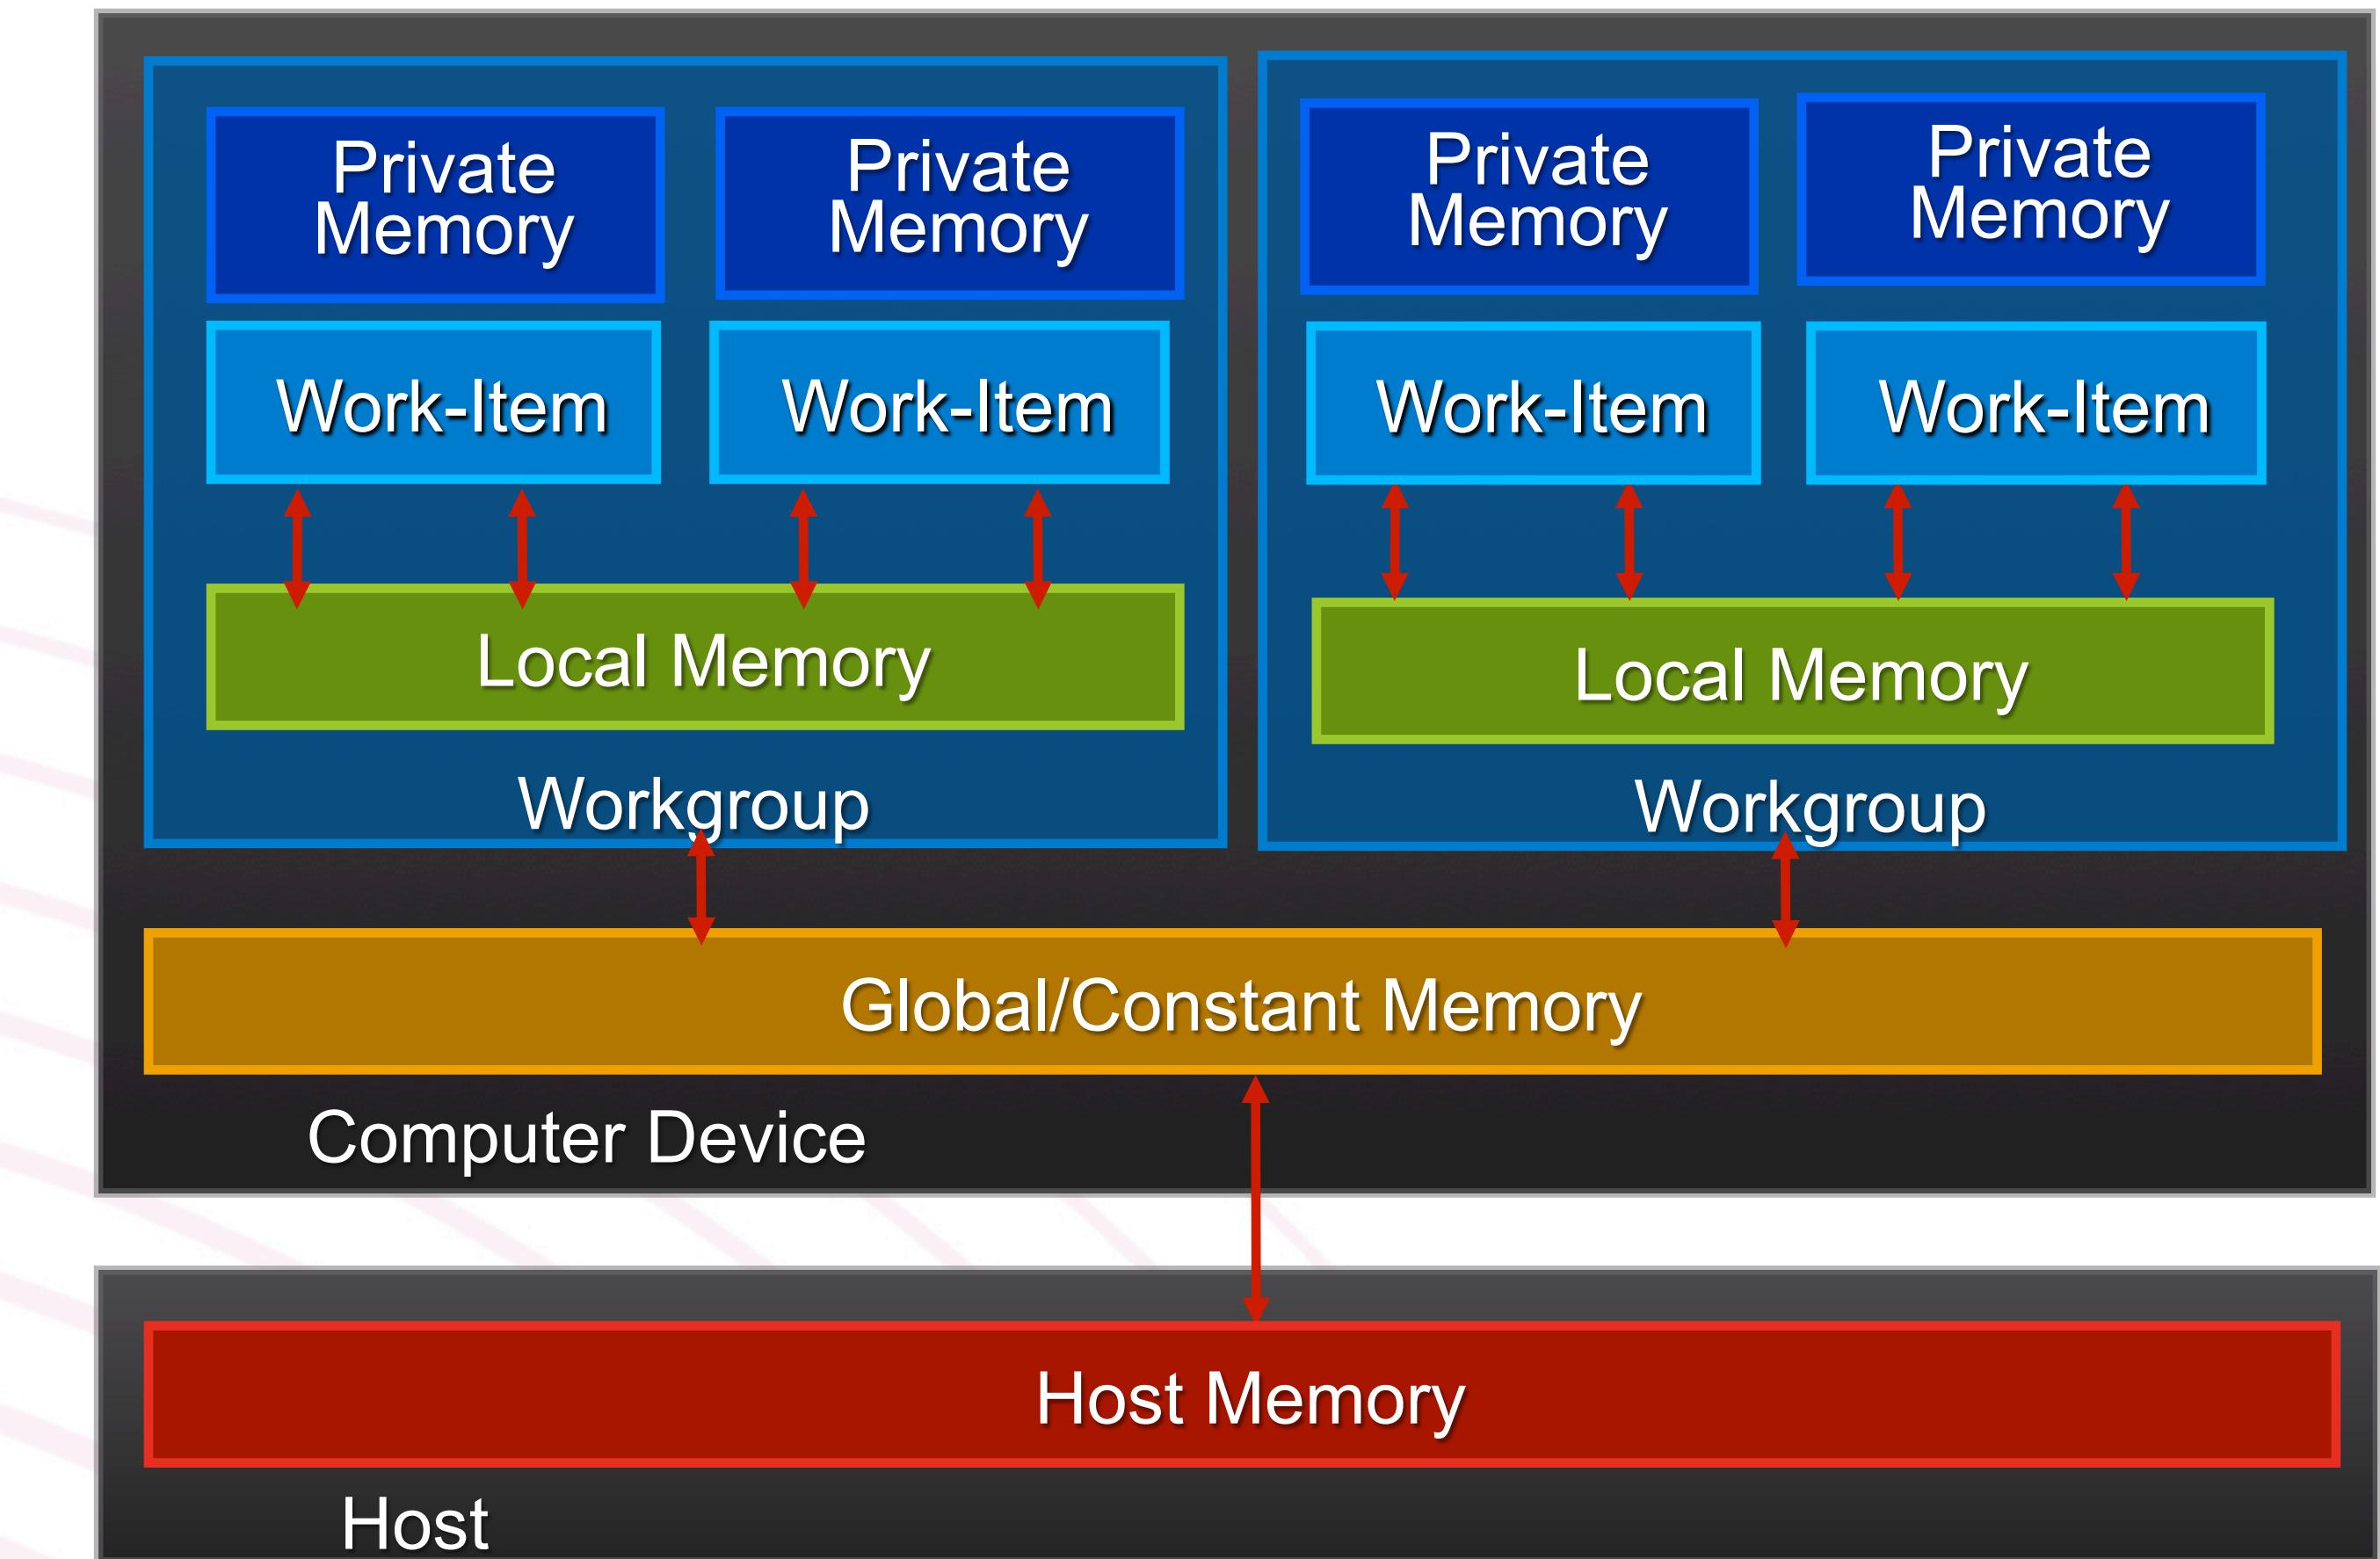
- Shared within a workgroup

- **Local Global/Constant Memory**

- Visible to all workgroups

- **Host Memory**

- On the CPU



- Memory management is explicit

You must move data from host -> global -> local and back

# OpenCL C Language

- **Derived from ISO C99**
  - No standard C99 headers, function pointers, recursion, variable length arrays, and bit fields
- **Additions to the language for parallelism**
  - Work-items and workgroups
  - Vector types
  - Synchronization
- **Address space qualifiers**
- **Optimized image access**
- **Built-in functions**

# OpenCL 1.1 - API

- Thread-safety
  - All API calls, except **clSetKernelArg**, are thread safe
- Sub-buffer objects
  - Create an object that represents a specific region in a buffer object
  - Easy and efficient mechanism to distribute regions of a buffer object across multiple devices
  - OpenCL™ synchronization mechanism ensures modifications to sub-buffer object reflected in appropriate region of parent buffer object

# OpenCL 1.1 - API

- User Events
  - **clEnqueue\*\*\*** commands can wait on event
  - In OpenCL™ 1.0, events can only refer to OpenCL™ commands
  - Need ability to enqueue commands that wait on an external, user defined, event
- Event CallBacks
  - **clSetEventCallbackFn** to register a user callback function
    - called when command identified by event has completed
    - Allows applications to enqueue new OpenCL™ commands based on event state changes in a non-blocking manner

# OpenCL 1.1 - Language

- **Implicit Conversions**

- OpenCL™ 1.0 requires widening for arithmetic operators

```
float4 a, b;
```

```
float c;
```

```
b = a + c; // c is widened to a float4 vector
```

```
// first and then the add is performed
```

- OpenCL™ 1.1 extends this feature for all operators
  - relational, equality, bitwise, logical, ternary

# OpenCL 1.1 - Language

- **3-component vector data types**
  - And everyone applauds....well almost everyone
- **cl\_khr\_byte\_addressable\_as\_core\_feature**
- **Atomic extensions are now core features**
  - **cl\_khr\_global\_int32\_{base | extended}\_atomics**
  - **cl\_khr\_local\_int32\_{base | extended}\_atomics**

# OpenCL 1.1 - Language

- New built-in functions
  - **get\_global\_offset**
  - clamp for integer data types
  - **async\_work\_group\_strided\_copy**
  - strided async copy of data from global <---> local memory
  - **shuffle** - construct a permutation of elements from 1 or 2 input vectors and a mask

# OpenCL 1.1 – OpenCL/OpenGL Sharing

- Improve performance of OpenCL/ OpenGL interoperability
  - Portable OpenCL / OpenGL sharing requires
    - a **glFinish** before **clEnqueueAcquireGLObjec**ts
    - a **clFinish** after **clEnqueueReleaseGLObjec**ts
  - **glFinish** / **clFinish** are heavyweight APIs

# OpenCL 1.1 – OpenCL/OpenGL Sharing

- Improve performance of OpenCL/ OpenGL interoperability
  - Create a OpenCL event from an OpenGL sync object
  - Create a OpenGL sync object from a OpenCL event
  - Allows for a finer grained waiting mechanism
    - Use **event\_wait\_list** argument for events that refer to OpenGL commands to complete
    - Use OpenGL sync APIs to wait for specific OpenCL™ commands to complete